

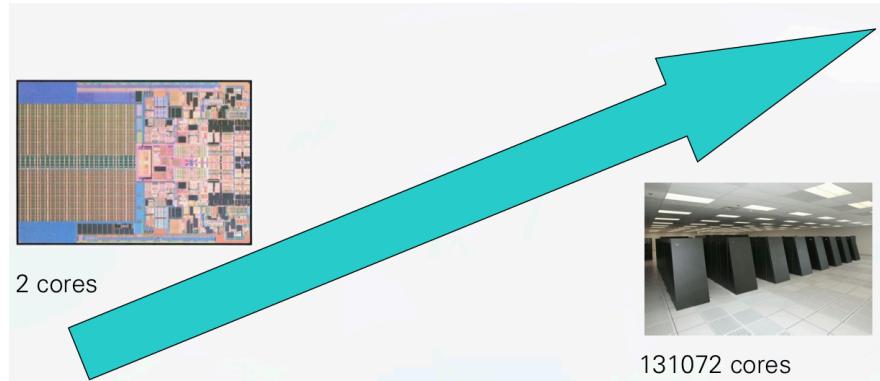
**ECE-451/ECE-566 - Introduction to
Parallel and Distributed Programming**

Performance Evaluation

Department of Electrical &
Computer Engineering
Rutgers University

**Performance Evaluation
Tools**

Motivation

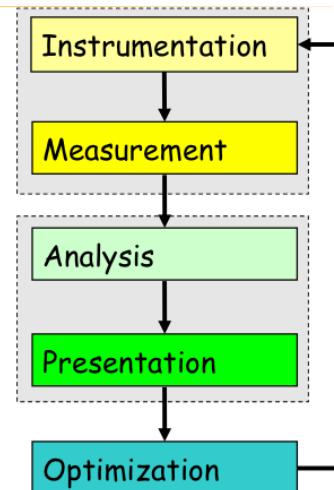


Scalability!!

Performance Optimization

- Expose factors
- Collect performance data
- Calculate metrics
- Analyze results
- Visualize results
- Identify problems
- Tune Performance

This is an “art”



Steps of Performance Evaluation

- Collect basic routine-level timing profile to determine where most time is being spent
- Collect routine-level hardware counter data to determine types of performance problems
- Collect callpath profiles to determine sequence of events causing performance problems
- Conduct finer-grained profiling and/or tracing to pinpoint performance bottlenecks
 - Loop-level profiling with hardware counters
 - Tracing of communication operations

Parallel Performance Properties

- Parallel code performance is influenced by both sequential and parallel factors?
- Sequential factors
 - Computation and memory use
 - Input / output
- Parallel factors
 - Thread / process interactions
 - Communication and synchronization

Performance Analysis Questions

- How does performance vary with different compilers?
- Is poor performance correlated with certain OS features?
- Has a recent change caused unanticipated performance?
- How does performance vary with MPI variants?
- Why is one application version faster than another?
- What is the reason for the observed scaling behavior?
- Did two runs exhibit similar performance?
- How are performance data related to application events?
- Which machines will run my code the fastest and why?
- Which benchmarks predict my code performance best?

Performance Considerations and Strategies

- The most important goal of performance tuning is to reduce a program's wall clock execution time.
 - Reducing resource usage in other areas, such as memory or disk requirements, may also be a tuning goal.
- Performance tuning is an iterative process used to optimize the efficiency of a program. It usually involves finding your program's *hot spots* and eliminating the *bottlenecks* in them.
 - *Hot Spot*: An area of code within the program that uses a disproportionately high amount of processor time
 - *Bottleneck*: An area of code within the program that uses processor resources inefficiently and therefore causes unnecessary delays
- Performance tuning usually involves profiling - using software tools to measure a program's run-time characteristics and resource utilization
- Use profiling tools and techniques to learn which areas of your code offer the greatest potential performance increase BEFORE you start the tuning process.
 - Then, target the most time consuming and frequently executed portions of a program for optimization.

An example

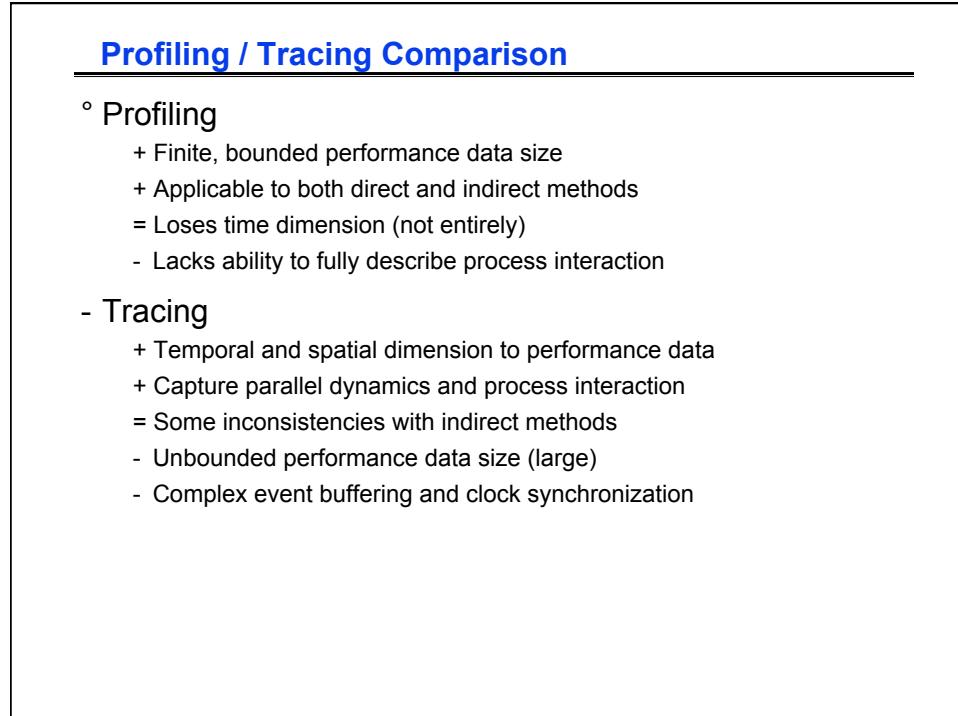
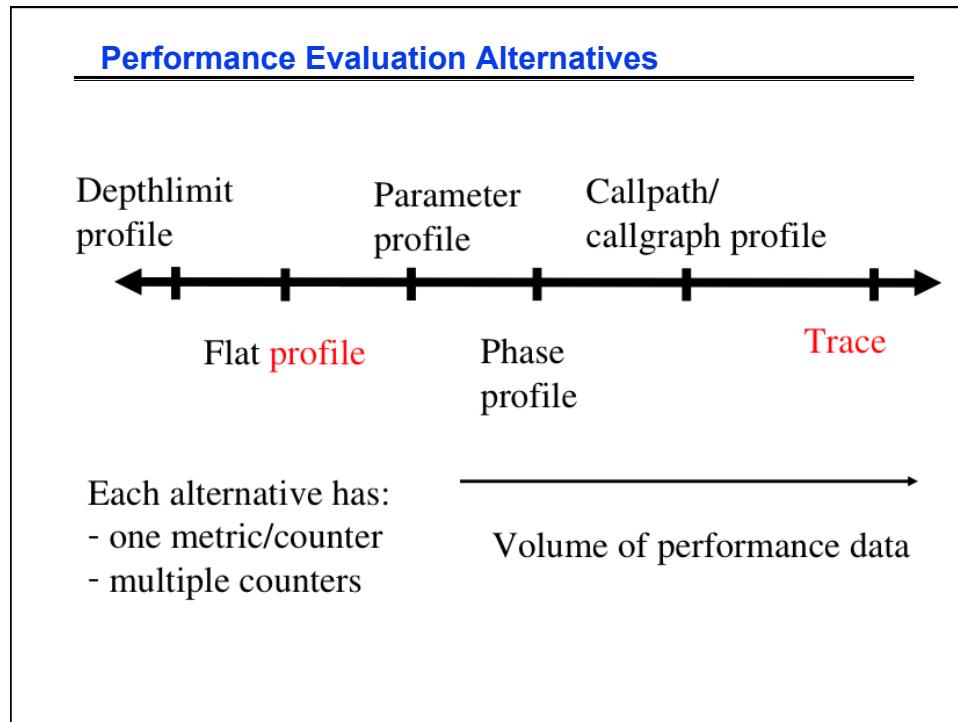
- Consider optimizing your underlying algorithm: an extremely fine-tuned O(N * N) sorting algorithm may perform significantly worse than a untuned O(N log N) algorithm
- For data dependent computations, benchmark based on a variety of realistic (both size and values) input data sets. Maintain consistent input data during the fine-tuning process.
- Take advantage of compiler and preprocessor optimizations when possible
- Finally, know when to stop - there are diminishing returns in successive optimizations. Consider a program with the following breakdown of execution time percentages for the associated parts of the program:

Procedure	%CPU Time
main()	13%
procedure1()	17%
procedure2()	20%
procedure3()	50%

- A 20% increase in the performance of procedure3() results in a 10% performance increase overall.
- A 20% increase in the performance of main() results in only a 2.6% performance increase overall.

Critical Issues

- Accuracy
 - Timing and counting accuracy depends on resolution
 - Any performance measurement generates overhead
 - Execution on performance measurement code
 - Measurement overhead can lead to intrusion
 - Intrusion can cause perturbation
 - Alters program behavior
- Granularity
 - How many measurements are made
 - How much overhead per measurement
- Tradeoff (general wisdom)
 - Accuracy is inversely correlated with granularity



Timers (example: IBM SP2 System)

Timer	Usage	Wallclock / CPU Time	Resolution	Languages	Portable?
time	shell / script	both	1/100th second	any	yes
timex	shell / script	both	1/100th second	any	yes
gettimeofday	subroutine	wallclock	microsecond	C/C++	yes
read_real_time	subroutine	wallclock	nanosecond	C/C++	no
rtc	subroutine	wallclock	microsecond	Fortran	no
irtc	subroutine	wallclock	nanosecond	Fortran	no
dtime_	subroutine	CPU	1/100th second	Fortran	no
etime_	subroutine	CPU	1/100th second	Fortran	no
mclock	subroutine	CPU	1/100th second	Fortran	no
timef	subroutine	wallclock	millisecond	Fortran	no
MPI_Wtime	subroutine	wallclock	microsecond	C/C++, Fortran	yes
AIX Trace Facility	shell / script / subroutine	wallclock	microsecond	any	no

Time

- The time command returns the total execution time of your program.
- The format of the output is different for the Korn shell and the C shell. The basic information is : Real time: the total wall clock (start to finish) time your program took to load, execute, and exit.
- User time: the total amount of CPU time your program took to execute.
- System time: the amount of CPU time spent on operating system calls in executing your program.
- The system and user times are defined differently across different computer architectures.
- Example csh time output:

```

 1 2 3 4 5 6 7 8
1.150u 0.020s 0:01.76 66.4% 15+3981k 24+10io 0pf+0w

```

- Explanation:
- 1.15 seconds of user CPU time
- 0.02 seconds of system (kernel) time used on behalf of user
- 1.76 seconds real time (wall clock time)
- 66.4% total CPU time (user+system) during execution as a percentage of elapsed time.
- 15 Kbytes of shared memory usage and 3981 Kbytes of unshared data space
- 24 block input operations and 10 block output operations
- no page faults
- no swaps

Time (cont.)

- ° Example ksh time output:

```
real 0m2.58s
user 0m1.14s
sys 0m0.03s
```

Explanation:

- 0 minutes, 2.58 seconds of wall clock time
- 0 minutes, 1.14 seconds of user CPU time
- 0 minutes, 0.03 seconds of system CPU time

gettimeofday()

- ° The gettimeofday() routine is part of the Standard C Library (libc.a) on most Unix systems. It returns the time in seconds and microseconds since midnight, January 1, 1970.
- ° Can be inserted anywhere within a C program and used to determine the start and end times of code fragments.
- ° Fortran programs must perform an interlanguage call in order to use the gettimeofday function.
- ° Timer resolution is hardware dependent. Microsecond resolution on the IBM SP systems.

```
#include <stdio.h>
#include <sys/time.h>
#include <time.h>

struct timeval start_time, end_time;

main()
{
    int total_usecs;
    gettimeofday(&start_time, (struct timeval*)0);
    /* do some work */
    gettimeofday(&end_time, (struct timeval*)0);
    total_usecs = (end_time.tv_sec-start_time.tv_sec) * 1000000 +
                  (end_time.tv_usec-start_time.tv_usec);
    printf("Total time was %d uSec.\n", total_usecs);
}
```

read_real_time() & time_base_to_time()

- Both of these routines are part of the Standard C Library (libc.a) on IBM AIX systems.
- Designed to be used for making high-resolution measurement of elapsed time, using the processor real time clock or time base registers.
- Nanosecond resolution on the IBM SP.
- The time_base_to_time() routine is used to guarantee correct time units across different IBM RS/6000 architectures.

```
#include <stdio.h>
#include <sys/time.h>

int main(void)
{
    timebasestruct_t start, finish;
    int secs, n_secs;

    read_real_time(&start, TIMEBASE_SZ);
    /* do some work */

    read_real_time(&finish, TIMEBASE_SZ);
    /* Make sure both values are in seconds and nanoseconds */
    time_base_to_time(&start, TIMEBASE_SZ);
    time_base_to_time(&finish, TIMEBASE_SZ);

    /* Subtract the starting time from the ending time */
    secs = finish.tb_high - start.tb_high;
    n_secs = finish.tb_low - start.tb_low;

    /* Fix carry from low-order to high-order during the measurement */
    if (n_secs < 0) {
        secs--;
        n_secs += 1000000000;
    }

    (void) printf("Sample time was %d seconds %d nanoseconds\n",
                 secs, n_secs);
    exit(0);
}
```

AIX Trace Facility

- AIX's built in Trace Facility provides the ability to timestamp many different system events.
- It can be used at both the command line level and subroutine level.
- Using the Trace Facility to collect timing statistics for an event is a three step process:
 - Determine the code numbers for the events you wish to trace by consulting the /usr/include/sys/trchkid.h file.
 - Generate a tracefile from either the command line or from within your program.
 - Produce a "readable" trace report from your tracefile after it has been completed.
- Example:
 - Get event tags from /usr/include/sys/trchkid.h for desired events. Thread create=465, thread terminate=467, wait lock=46d.
 - Remove any existing tracefile (can cause problems): rm trcfile
 - Execute program, creating a tracefile with specified events: trace -a -o trcfile -j 465,467,46d; dotprod; trcstop
 - Produce a readable trace report: trcrpt -o trace.report -O pid=on trcfile
 - View report using your favorite editor.

AIX Trace Facility

```

Fri Jul 19 00:02:05 2002
System: AIX frost067 Node: 5
Machine: 006008594C00
Internet Address: 86092F43 134.9.47.67
The system contains 16 cpus, of which 16 were traced.
Buffering: Kernel Heap
This is from a 32-bit kernel.
Tracing only these hooks, 465,467,46d

trace -a -o trcfile -j 465,467,46d

ID PID I ELAPSED_SEC DELTA_MSEC APPL SYSCALL KERNEL INTERRUPT
001 70640 0.000000000 0.000000 TRACE ON channel
0
Fri Jul 19 00:02:
05 2002
465 -1 0.000482511 0.482511 thread_create: pid=7825
0 tid=226787 priority=60 policy=0 wait_on_lock: pid=85616
46D 85616 0.000495484 0.012973
tid=238427 lockaddr=5819F0
465 -1 0.027897260 27.401776 thread_create: pid=7731
8 tid=146769 priority=61 policy=0 wait_on_lock: pid=7731
467 77318 0.028225717 0.328457 thread_terminate: pid=7
7318 tid=146769
465 -1 0.028235827 0.010110 thread_create: pid=7731
8 tid=206215 priority=61 policy=0 wait_on_lock: pid=77318
46D -1 0.028251016 0.015189
tid=238429 lockaddr=5819F0
465 -1 0.028633968 0.382952 thread_create: pid=7731
8 tid=151151 priority=61 policy=0 wait_on_lock: pid=7731
467 77318 0.028795141 0.161173 thread_terminate: pid=7
7318 tid=206215
467 77318 0.028861449 0.066308 thread_terminate: pid=7
7318 tid=151151
465 -1 0.028876390 0.014941 thread_create: pid=7731
8 tid=233391 priority=61 policy=0 wait_on_lock: pid=77318
46D -1 0.028886636 0.010246
tid=238429 lockaddr=5819F0
467 77318 0.029303598 0.416962 thread_terminate: pid=7
7318 tid=233391
002 -1 0.055766970 26.463372 TRACE OFF channel
0000 Fri Jul 19 00:02:05 2002

```

Profilers - prof

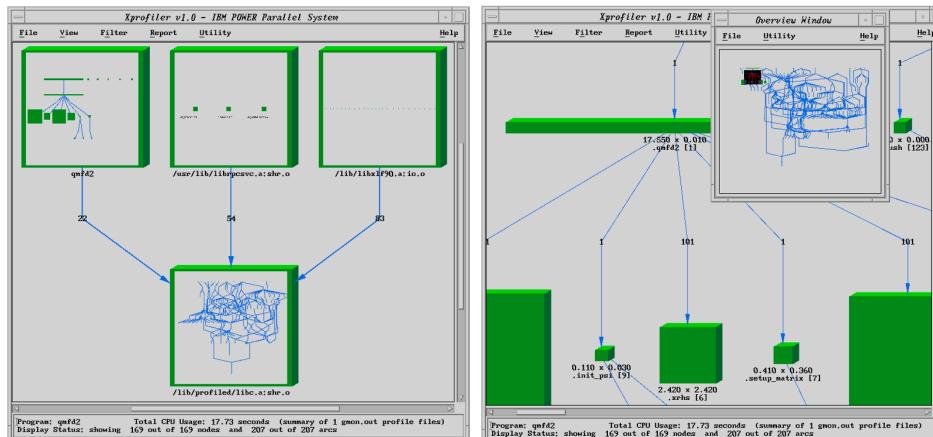
- The prof utility is included in most Unix systems. It is used to profile program execution at the procedure level.
- IBM has enabled the mpcc and mpxl compiler commands to use prof for parallel programs.
- prof displays the following information: The name of each procedure - in descending order of processing activity
- The percentage of the program's CPU time used by each procedure.
- The execution time in seconds for all references by each procedure.
- The number of times the procedure was called.
- The average time in milliseconds for a call to each procedure.

Name	%Time	Seconds	Cumsecs	#Calls	msec/call
.fft	51.8	0.59	0.59	1024	0.576
.main	40.4	0.46	1.05	1	460
.bit_reverse	7.9	0.09	1.14	1024	0.088
.cos	0.0	0.00	1.14	256	0.00
.sin	0.0	0.00	1.14	256	0.00
.catopen	0.0	0.00	1.14	1	0.
.setlocale	0.0	0.00	1.14	1	0.
.doprnt	0.0	0.00	1.14	7	0.
.flsbuf	0.0	0.00	1.14	11	0.0
.xflsbuf	0.0	0.00	1.14	7	0.
._wrchk	0.0	0.00	1.14	1	0.
._findbuf	0.0	0.00	1.14	1	0.
._xwrite	0.0	0.00	1.14	7	0.
.free	0.0	0.00	1.14	2	0.
.free_y	0.0	0.00	1.14	2	0.
.write	0.0	0.00	1.14	7	0.
.exit	0.0	0.00	1.14	1	0.
.memchr	0.0	0.00	1.14	19	0.
.atoi	0.0	0.00	1.14	1	0.
._nl_langinfo_std	0.0	0.00	1.14	4	0.
.gettimeofday	0.0	0.00	1.14	8	0.
.printf	0.0	0.00	1.14	7	0.

xprofiler

- xprofiler is an X Windows based profiling tool distributed with IBM's Parallel Environment software. It is based upon the gprof profiling utility, however it provides a graphical representation of profile data in addition to all of the usual gprof reports.
- Graphical information is organized into 3 main components:
 - [Library Cluster Boxes](#)
 - [Function Boxes](#)
 - [Call Arcs](#)
- xprofiler can profile at both the subroutine level, and at the source statement level. Profiling also includes any calls made to library functions.
- Time sampled profiling:
 - Like many profiling tools, xprofiler acquires its data by keeping track of the executing program's location whenever it is interrupted.
 - Interrupts occur at set intervals and are commonly called "ticks". Under AIX, a tick occurs every 1/100th second.
 - Those portions of the program which accumulate the most "ticks" reflect the areas where the program spends most of its time.
- Filters, zooming, and other options allow you to limit displays to only those portions of the call tree you are interested in analyzing.

Xprofile (cont.)



Hardware Performance Counters

- For many years, hardware engineers have designed in specialized registers to measure the performance of various aspects of a microprocessor
- HW performance counters provide application developers with valuable information about code sections that can be improved
- Hardware performance counters can provide insight into:
 - Whole program timing
 - Cache behaviors
 - Branch behaviors
 - Memory and resource contention and access patterns
 - Pipeline stalls
 - Floating point efficiency
 - Instructions per cycle
 - Subroutine resolution
 - Process or thread attribution

What's PAPI?

- Middleware that provides a consistent and efficient programming interface for the performance counter hardware found in most major microprocessors.
- Started as a Parallel Tools Consortium project in 1998
 - Goal was to produce a specification for a **portable** interface to the hardware performance counters
- Countable events are defined in two ways:
 - Platform-neutral **Preset Events** (e.g., PAPI_TOT_INS)
 - Platform-dependent **Native Events** (e.g., L3_CACHE_MISS)
- Preset Events can be **derived** from multiple Native Events (e.g., PAPI_L1_TCM might be the sum of L1 Data Misses and L1 instruction Misses on a given platform)

PAPI Hardware Events

° Preset Events

- Standard set of over 100 events for application performance tuning
- No standardization of the exact definition
- Mapped to either single or linear combinations of native events on each platform
- Use **papi_avail** to see what preset events are available on a given platform

° Native Events

- Any event countable by the CPU
- Same interface as for preset events
- Use **papi_native_avail** utility to see all available native events

° Use **papi_event_chooser** utility to select a compatible set of events

PAPI Counter Interfaces

° PAPI provides 3 interfaces to the underlying counter hardware:

* A Low Level API manages hardware events (present and native) in user defined groups called *EventSets*. Meant for experienced application programmers wanting fine-grained measurements.

* A High Level API provides the ability to start, stop and read the counters for a specified list of events (preset only). Meant for programmers wanting simple event measurements.

3rd Party and GUI Tools



PAPI PORTABLE LAYER

PAPI HARDWARE SPECIFIC LAYER

Kernel Extension

Operating System

Perf Counter Hardware

PAPI High Level Calls

```

PAPI_num_counters()
    ♦   get the number of hardware counters available on the system
PAPI_flips (float *rtime, float *ptime, long long *flips, float *mflips)
    ♦   simplified call to get Mflips/s (floating point instruction rate), real and processor time
PAPI_flops (float *rtime, float *ptime, long long *flops, float *mflops)
    ♦   simplified call to get Mflops/s (floating point operation rate), real and processor time
PAPI_ipc (float *rtime, float *ptime, long long *ins, float *ipc)
    ♦   gets instructions per cycle, real and processor time
PAPI_accum_counters (long long *values, int array_len)
    ♦   add current counts to array and reset counters
PAPI_read_counters (long long *values, int array_len)
    ♦   copy current counts to array and reset counters
PAPI_start_counters (int *events, int array_len)
    ♦   start counting hardware events
PAPI_stop_counters (long long *values, int array_len)
    ♦   stop counters and return current counts

```

Example: Low Level API

```

#include "papi.h"
#define NUM_EVENTS 2
int Events[NUM_EVENTS]={PAPI_FP_OPS,PAPI_TOT_CYC},
int EventSet;
long long values[NUM_EVENTS];

/* Initialize the Library */
retval = PAPI_library_init (PAPI_VER_CURRENT);
/* Allocate space for the new eventset and do setup */
retval = PAPI_create_eventset (&EventSet);
/* Add Flops and total cycles to the eventset */
retval = PAPI_add_events (&EventSet,Events,NUM_EVENTS);

/* Start the counters */
retval = PAPI_start (EventSet);

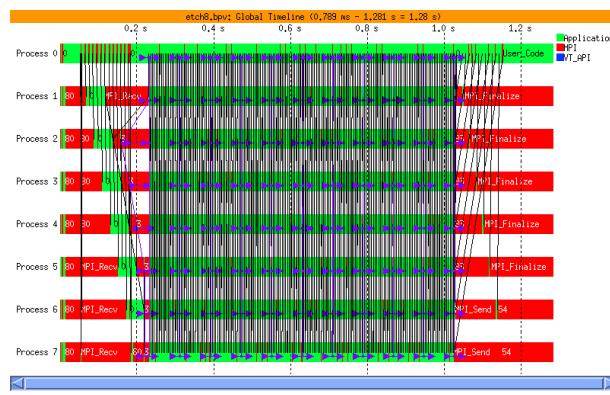
do_work(); /* What we want to monitor*/

/*Stop counters and store results in values */
retval = PAPI_stop (EventSet,values);

```

Tracing tools

- The parallel program is monitored while it is executed
- Monitoring produces performance data (TRACE) that is interpreted in order to reveal areas of poor performance.
- Tracing tools examples:
 - Vampir
 - Paraver



Vampir (Visualization and Analysis of MPI Resources)

- VAMPIR 2.0 is a post-mortem trace visualization tool from Pallas GmbH

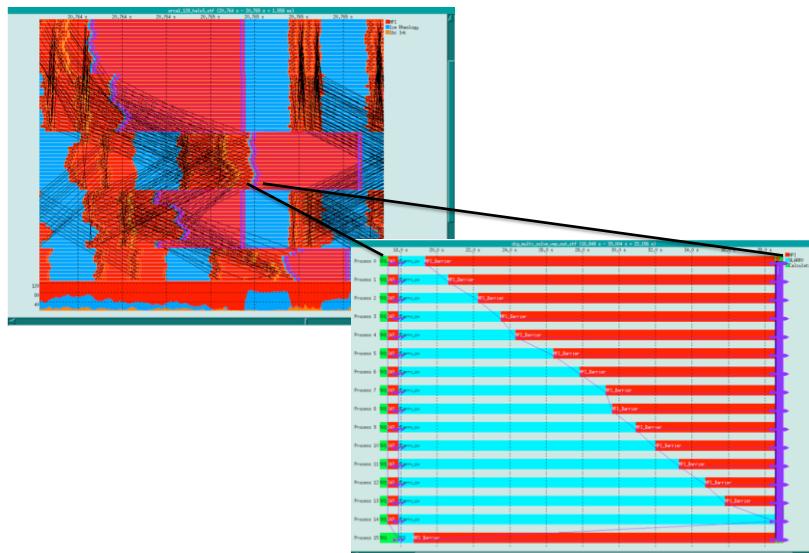
<http://www.pallas.com>

It uses the profile extensions to MPI and permits analysis of the message events where data is transmitted between processors during execution of a parallel program. It has a convenient user-interface and an excellent zooming and filtering. Global displays show all selected processes.

Vampir

- *Global Timeline*: detailed application execution over time axis
- *Activity Chart*: presents per-process profiling information
- *Summary Chart*: aggregated profiling information
- *Communication Statistics*: message statistics for each process pair
- *Global Communication Statistics*: collective operations statistics
- *I/O Statistics*: MPI I/O operation statistics
- *Calling Tree*: global dynamic calling tree

Vampir



Paraver

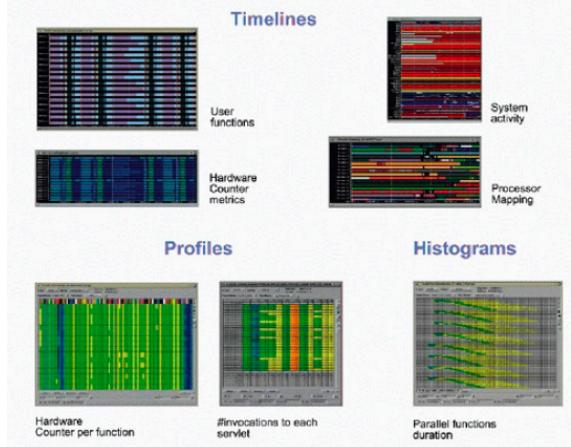
- Powerful flexible parallel program visualization tool based on an easy-to-use Motif GUI (graphical user interface)
 - Current version under Linux and Windows (wxWidgets). LGPL license
- Developed by :

European Center for Parallelism of Barcelona (CEPBA) /
Barcelona Supercomputing Center
(Universitat Politecnica de Catalunya)

<http://www.bsc.es/>

Paraver

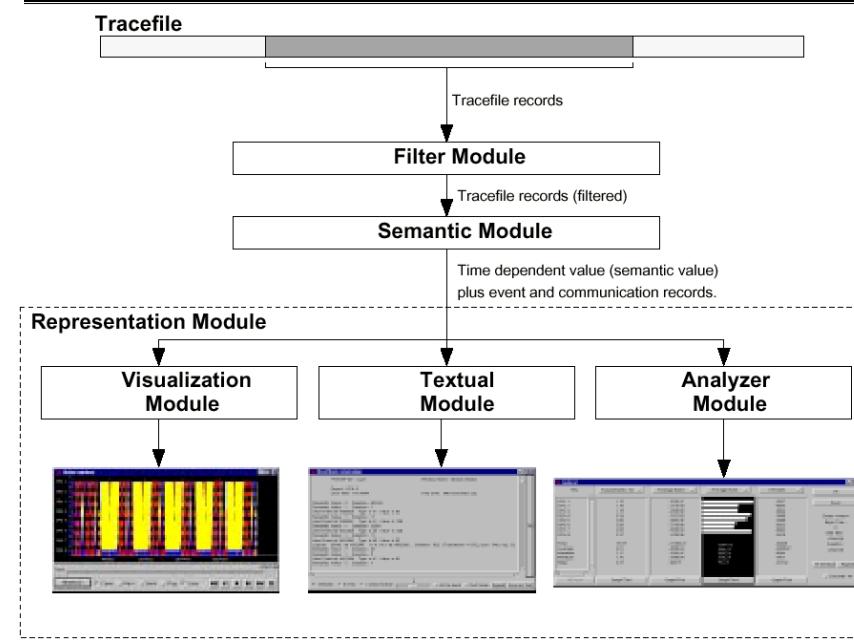
- Single Visualization Environment for MPI, OpenMP, Mixed, MLP, Java, OS activity
- Multiple platforms: IRIX, AIX, LINUX, Tru64

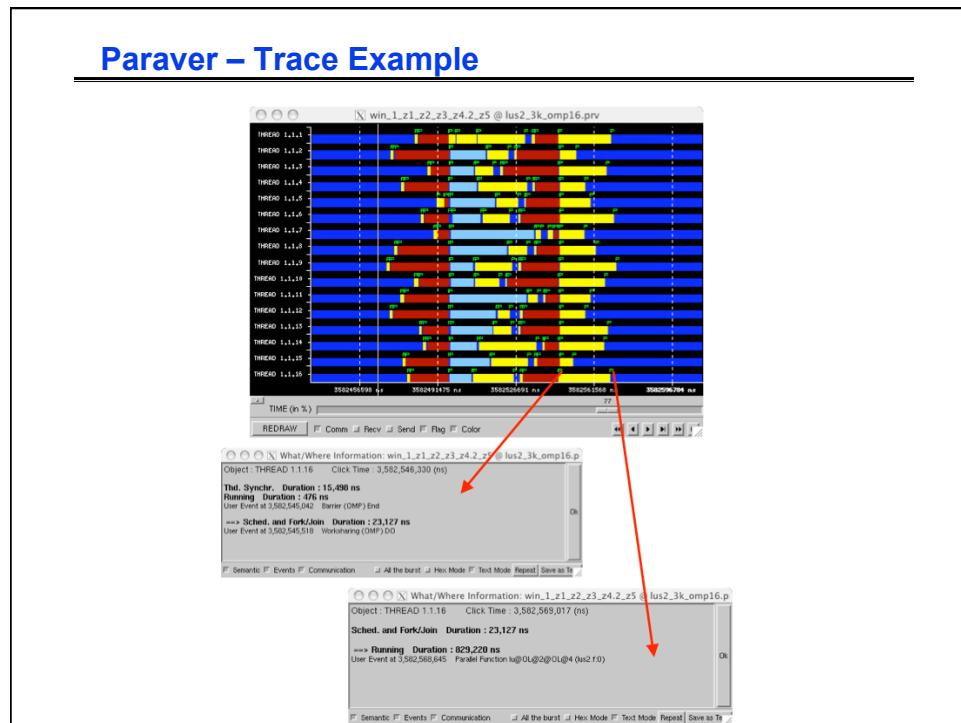
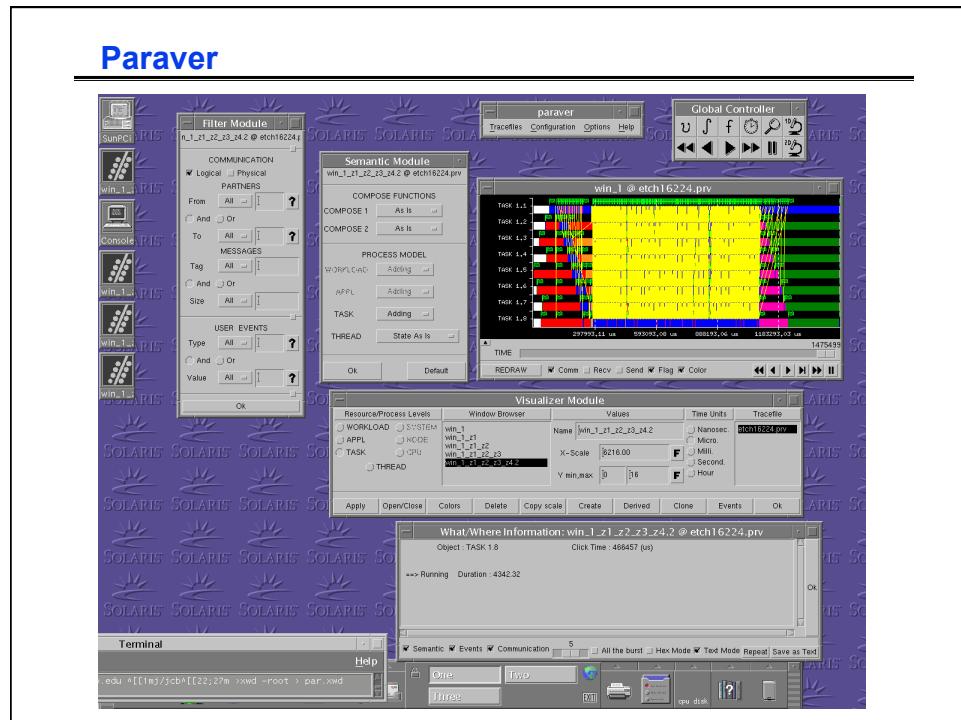


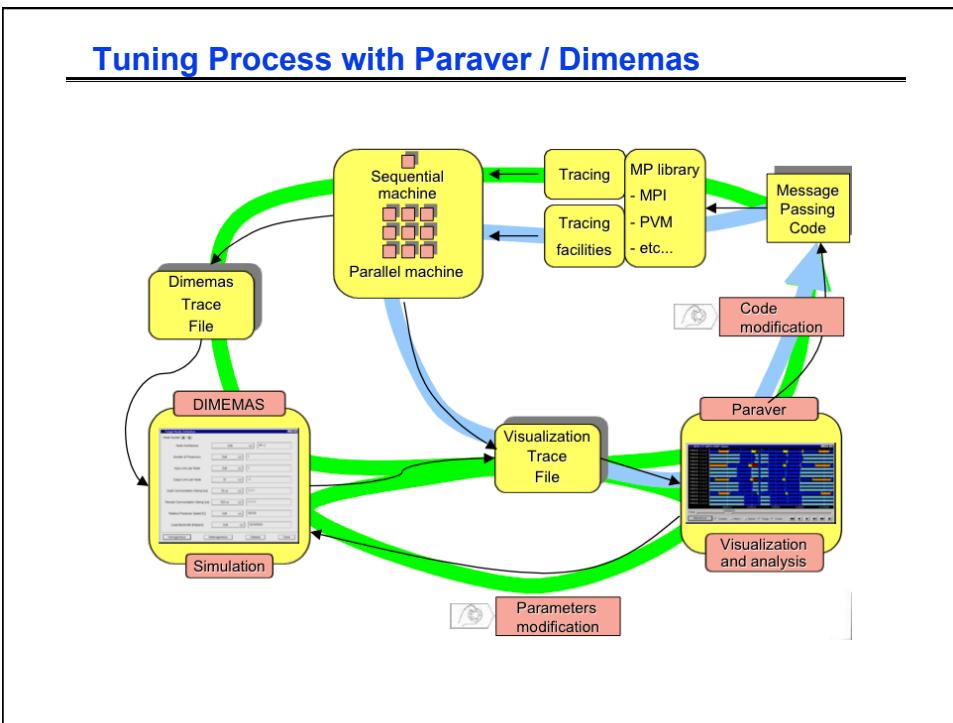
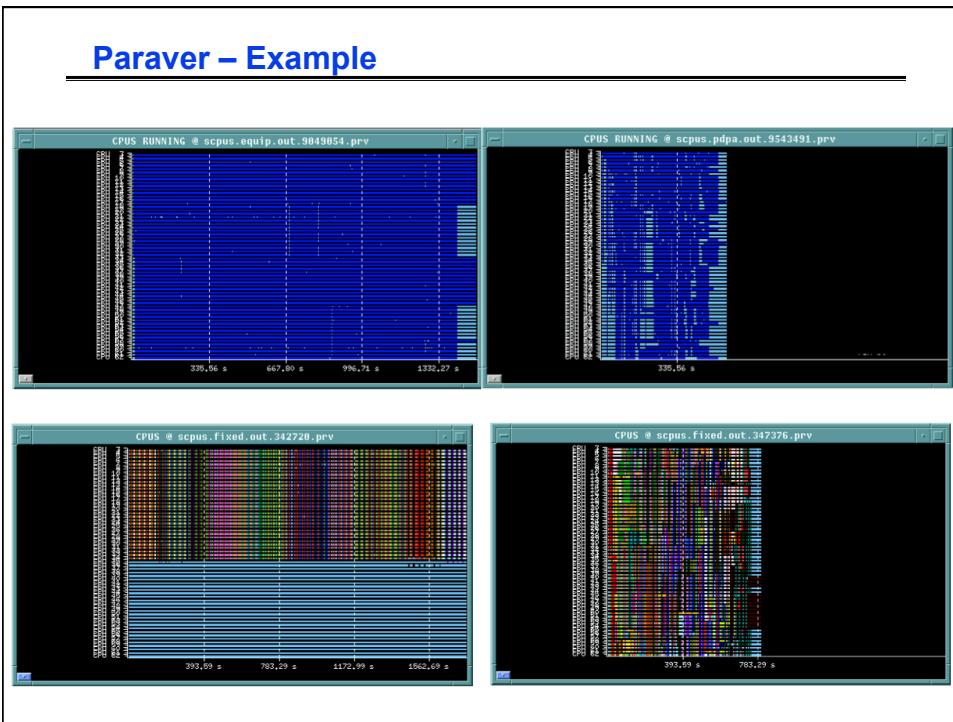
Paraver

- Paraver is designed to visualize and analyze:
 - Communication and load balance
 - Combining OpenMP and MPI
 - Hardware performance and counters
- Usage
 - Compile programs with special libraries
 - Run programs to produce trace files
 - View and analyze traces
 - Designed to help in program understanding and optimization

Paraver







Vtune

- Presentation / Discussion in class...

TAU Performance System Project

- Tuning and Analysis Utilities (15+ year project effort)
- Performance system framework for HPC systems
 - Integrated, scalable, and flexible
 - Target parallel programming paradigms
- Integrated toolkit for performance problem solving
 - Instrumentation, measurement, analysis, and visualization
 - Portable performance profiling and tracing facility
 - Performance data management and data mining
- Partners
 - LLNL, ANL, LANL
 - Research Center Jülich, TU Dresden

What is TAU?

- TAU is a performance evaluation tool
- It supports parallel profiling and tracing
- Profiling shows you how much (total) time was spent in each routine
- Tracing shows you when the events take place in each process along a timeline
- TAU uses a package called PDT for automatic instrumentation of source code
- Profiling and tracing can measure time as well as hardware performance counters from your CPU
- TAU can automatically instrument your source code (routines, loops, I/O, memory, phases, etc.)
- TAU runs on all HPC platforms and it is free (BSD style license)
- TAU has instrumentation, measurement and analysis tools
 - Paraprof is TAU's 3D profile browser
- To use TAU, you need to set a couple of environment variables and substitute the name of your compiler with a TAU shell script

TAU Performance System Interfaces

- PDT [U. Oregon, LANL, FZJ] for instrumentation of C++, C99, F95 source code
- PAPI [UTK] for accessing hardware performance counters data
- DyninstAPI [U. Maryland, U. Wisconsin] for runtime instrumentation
- KOJAK [FZJ, UTK]
 - Epilog trace generation library
 - CUBE callgraph visualizer
 - Opari OpenMP directive rewriting tool
- Vampir/VNG Trace Analyzer [TU Dresden]
- VTF3/OTF trace generation library [TU Dresden]
- Paraver trace visualizer [CEPBA/BSC]
- Jumpshot-4 trace visualizer [MPICH, ANL]
- JVMPPI from JDK for Java program instrumentation [Sun]
- Paraprof profile browser/PerfDMF database supports”
 - TAU format
 - Gprof [GNU]
 - Etc.

Using TAU: A brief Introduction

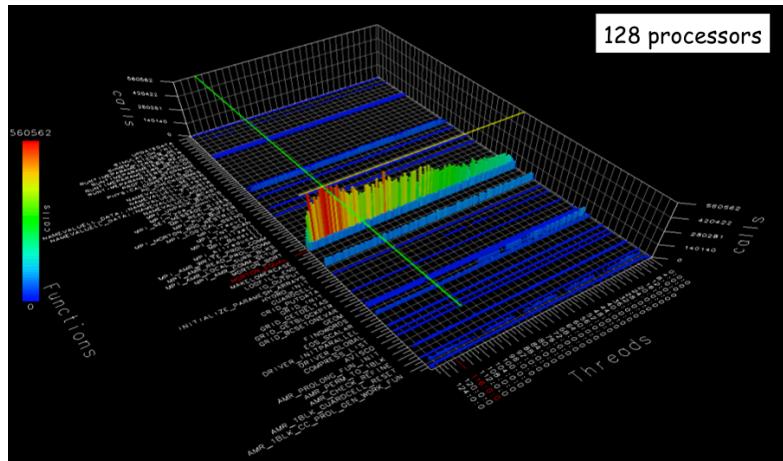
- TAU supports several measurement options (profiling, tracing, profiling with hardware counters, etc.)
- Each measurement configuration of TAU corresponds to a unique stub makefile that is generated when you configure it
- To instrument source code using PDT
 - Choose an appropriate TAU stub makefile in <arch>/lib:

```
% export TAU_MAKEFILE=/projects/tau/tau_latest/x86_64/lib/Makefile.tau-mpi-pdt
% export TAU_OPTIONS=-optVerbose ... (see tau_compiler.sh -help)
```

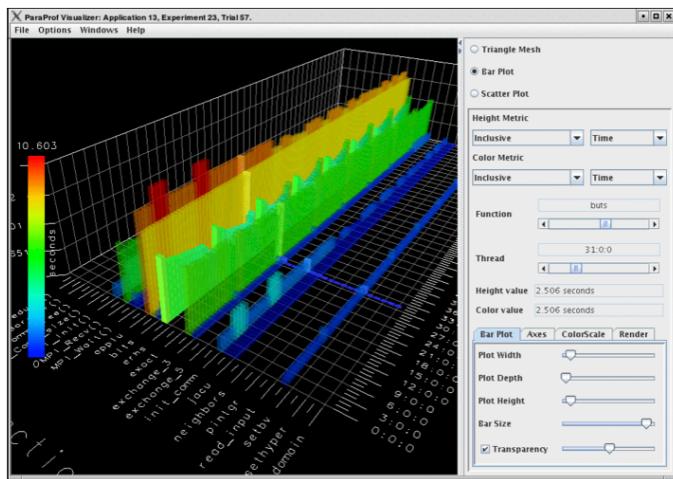
And use tau_f90.sh, tau_cxx.sh or tau_cc.sh as Fortran, C++ or C compilers:

```
% mpif90 foo.f90
changes to
% tau_f90.sh foo.f90
```
- Execute application and analyze performance data:
 - At runtime, if more than one metric is measured
 - export TAU_METRICS=TIME:PAPI_FP_INS:PAPI_NATIVE_<native_event_name>
 - Use papi_native_avail, papi_avail, and papi_event_chooser to select these preset and native event names
 - pprof (for text based profile display)
 - paraprof (for GUI)

ParaProf – 3D Full Profile Bar Plot



ParaProf Bar Plot (Zoom in/out +/-)



Vampir – Trace Zoomed



Benchmarks

Micro-benchmarks

- ° Memory: design micro-benchmarks to measure
 - Access time to memory of one processor
 - Maximum bandwidth
 - Effect of contention in the interconnection network
- ° Communication: design different versions of micro-benchmarks to measure the bandwidth of MPI

Micro-benchmarks

◦ Memory

```
for(i=0; i<N; i+=B)
    s+=a[i]
```

BW... max
B > cache line size

```
for(i=0; i<N; i+=B)
    p=p->next;
```

Back-to-back latency

```
for(i=0; i<N; i+=B){
    p=p->next;
    for(j=1/j<J;j++)
        s+=1;
}
```

True unloaded latency

HPC benchmarks

◦ Mbench (Intel Labs)

- *A Memory Consistency Benchmark for OpenMP*
- Applications (7)
 - Body tracker, cloth, face, fb_est, fimi, nee and ode
- Kernels (8)
 - adjust_velocity, binomial_tree, blackScholes, dual_intersection, fft, pcg, svd and svm
- Primitives (9)
 - dense_mvm, dense_mvm_sym, gups, sparse_mvm, sparse_mvm_sym, sparse_mvm_trans

◦ Parsec (Princeton university)

- *The Princeton Application Repository for Shared-Memory Computers*
- Applications (10)
 - Blackscholes, bodytrack, facesim, ferret, fluidanimate, freqmine, raytrace, swaptions, vips, x264
- Kernels (3)
 - Canneal, dedup, streamcluster

◦ NPB (NASA)

- NASA Parallel Benchmarks (different types and classes/sizes)

◦ Others

- SSCA (Graph Analysis)
- LAMMPS (Molecular dynamics – MPI App)
- Sequoia Bechmark (MPI Apps)
 - AMG , IRS, LAMMPS, Pynamic and UMT

HPC Challenge benchmark (7 tests)

1. [HPL](#) - the Linpack TPP benchmark which measures the floating point rate of execution for solving a linear system of equations.
2. [DGEMM](#) - measures the floating point rate of execution of double precision real matrix-matrix multiplication.
3. [STREAM](#) - a simple synthetic benchmark program that measures sustainable memory bandwidth (in GB/s) and the corresponding computation rate for simple vector kernel.
4. [PTRANS](#) (parallel matrix transpose) - exercises the communications where pairs of processors communicate with each other simultaneously. It is a useful test of the total communications capacity of the network.
5. [RandomAccess](#) - measures the rate of integer random updates of memory (GUPS).
6. [FFT](#) - measures the floating point rate of execution of double precision complex one-dimensional Discrete Fourier Transform (DFT).
7. Communication bandwidth and latency - a set of tests to measure latency and bandwidth of a number of simultaneous communication patterns; based on [b_eff](#) (effective bandwidth benchmark)